



Figure 1. RC-GRPO vs. existing GRPO post-training. Left: scoring candidates against raw frames (target-set mismatch) and broadcasting one sequence-level advantage mis-specify GRPO’s rewards R and advantages A . Right: (1) rank w.r.t. reachable reconstructions $\bar{s}'=D(Q(E(s')))$; (2) project updates to preserve raw first

st-order progress; (3) assign per-frame-block advantages from return-to-go. Final metrics remain vs. raw s' . Backbone, tokenizer, and decoder stay frozen.